



COMP3311 Database Management Systems
Spring 2022

+ SQL

- SQL is considered one of the major reasons for the success of relational model in the database industry
- It provides an industry wide standard for database access
 - In practice different product support different dialects of SQL
- It is a **declarative language** for users to specify what the result of the query should be, DBMS decides operations and order of execution
- SQL is designed for **data definition**, **data manipulation**, and **data control**, powerful enough to retrieve any piece of data from database

+ Three Types of SQL Statements

3

- Data definition language (DDL)
 - Statements to define the database **schema**
- Data manipulation language (DML)
 - Statements to manipulate the database **instances**
- Data control language (DCL)
 - Statements such as GRANT and REVOKE which mainly deal with the rights, permissions, and other controls of the database system

... SQL is used in all commercial RDBMSs
... different RDBMSs may have different syntax/features of SQL

+ History of SQL

- SEQUEL (Structured English Query Language) was first introduced from IBM Research for SYSTEM R in 1970s
- SQL:1986 introduced by ANSI and ISO
- SQL:1992 was a revised and much expanded version
- SQL:1999 extends SQL with object-oriented concepts
- SQL:2003 introduced XML features
- SQL:2016 adds JSON

+ SQL Outline

■ DDL

- Database Definition

■ DML

- Database Modification
- SELECT Statements

■ Database Views

■ More on Database Constraints

+ Data Definition Language (DDL)

The SQL DDL allows the specification of:

- The **schema** for each relation and their attributes
- The **types of values** associated with each attribute
 - char, varchar, int, smallint, numeric, real, double precision, float, date, time, timestamp...
- **User-defined types** and **domains**
- **Integrity constraints**
 - domain, key, foreign key, general
- The **physical storage structure** of each relation on disk
- The **set of indices** to be maintained for each relation

+ DDL - Data Definition Language

7

■ Basic SQL DDL Statements

- CREATE TABLE
- DROP TABLE
- ALTER TABLE
- CREATE DOMAIN

.. create/drop/modify **views** and **indexes** are also part of DDL, to be discussed later

... we don't discuss the specification of physical storage and transactions in this course

+ CREATE TABLE

- It creates a new relation, by specifying its **name**, **attributes** and **constraints**
 - The definitions are recorded the table definition in the **system catalog** (aka **data dictionary**)
 - The **key**, **entity** and **referential integrity constraints** are specified within the statement
 - The **domain constraint** is specified for each attribute by giving a valid (e.g., SQL99) data type and (optionally) excluding NULL from the domain
 - Valid data types include **INT**, **CHAR**, **DATE**, **DECIMAL**, etc.
 - Data type of an attribute can be specified directly or by declaring a domain (**CREATE DOMAIN**)

+ CREATE TABLE Syntax

9

CREATE TABLE <table name>

(<column name> <column type> [<attribute constraint>]

{, <column name> <column type> [<attribute constraint>] }

[<table constraint> {, <table constraint> }]);

Notations:

KEY WORDS, “,” “(” “)” “.”

<name>,

{repeat 0-n times},

[optional]

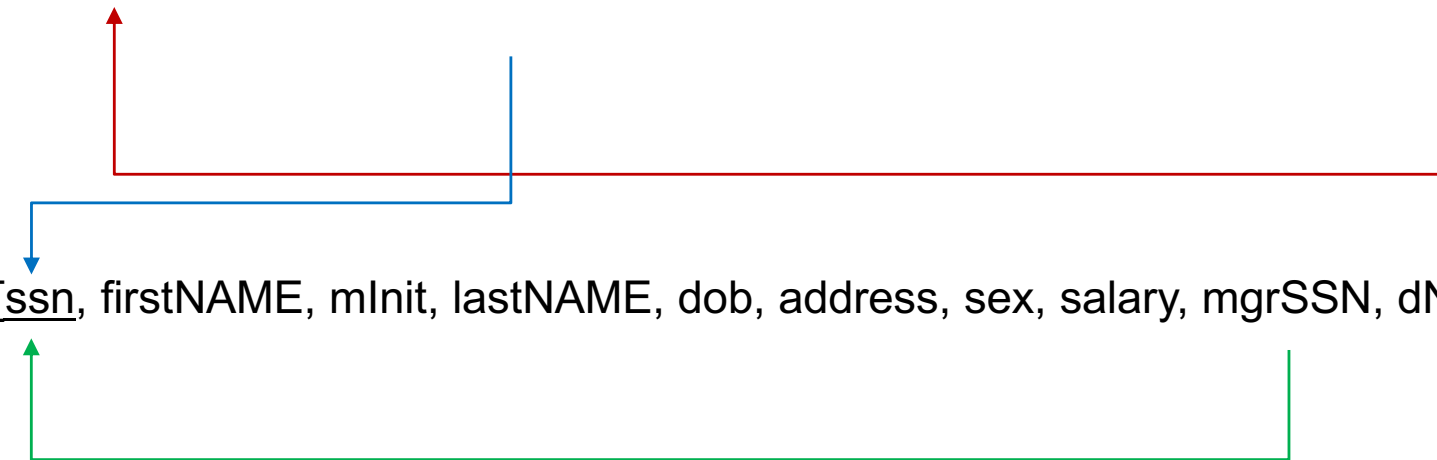
...you should keep an SQL syntax quick reference handy

+ CREATE TABLE Example

10

Department [dNumber, dName, mgrSSN, mgrStartDate]

Employee [ssn, firstNAME, mInit, lastNAME, dob, address, sex, salary, mgrSSN, dNum]



✎ The domain type of an attribute is enforced by the DBMS whenever tuples are added or modified, together with other constraints

... the domain type of each attribute needs to be specified (not captured in ER diagrams)

+ CREATE TABLE Example (1)

11

```
CREATE TABLE Employee
(  firstName      VARCHAR (15)      NOT NULL,
   mInit          CHAR,
   lastName       VARCHAR (15)      NOT NULL,
   ssn            CHAR (9)          NOT NULL,
   dob            DATE,
   address        VARCHAR (30),
   sex            CHAR,
   salary         DECIMAL (10, 2),
   mgrSSN         CHAR (9),
   dNum           INT              NOT NULL,
   PRIMARY KEY (ssn),
   FOREIGN KEY (mgrSSN) REFERENCES Employee (ssn),
   FOREIGN KEY (dNum) REFERENCES Department(dNumber) );
```

+ CREATE TABLE Example (2)

12

Constraints can be given a name:

```
CREATE TABLE Employee
(   firstName      VARCHAR (15)      NOT NULL,
    .....
    ssn             CHAR (9)          NOT NULL,
    mgrSSN          CHAR (9),
    dNum            INT               NOT NULL,
    CONSTRAINT empPK PRIMARY KEY (ssn),
    CONSTRAINT smpMgrFK FOREIGN KEY (mgrSSN)
        REFERENCES Employee (ssn),
    CONSTRAINT empDNumFK FOREIGN KEY (dNum)
        REFERENCES Department (dNumber)
);
```

+ Basic Types

13

char (n)	Fixed length character string with length n
varchar (n)	Variable-length character string with maximum length n
int	An integer (a finite subset of the integers that is machine-dependent)
smallint	A small integer (a machine-dependent subset of the integer domain type)
numeric (p, d)	A number with a total of p digits and d digits to the right of the decimal point
real	Floating-point and double-precision floating-point
double precision	Numbers with machine-dependent precision
float (n)	Floating point number, with user-specified precision of at least n digits
date	A date containing a (4 digit) year, month and day of month
time	The time of day, in hours, minutes and seconds
timestamp	A combination of date and time

...**Null** values are allowed in all the domain types

+ User-defined Types

14

■ Using CREATE TYPE clause

```
create type id_type as object (id numeric(10));
```

```
create type location_type as object (  
    address varchar(100),  
    country varchar(20)  
);
```

```
create table customer (  
    id id_type, // Oracle disallows a user-defined type as PK  
    add location_type  
);
```

👉 Not all relational systems support user-defined types

... this is a privileged operation
... complex user-defined types with operations/methods are possible
... object type inheritance is possible too

+ User-defined Domains

15

- The **CREATE DOMAIN** clause is used to define a new domain

```
create domain hourly_wage numeric(5,2);
```

- Differences between user-defined types and domains:
 - Domains can have constraints specified on them (such as **not null**) and can have default values defined for variables of the domain; types cannot.
 - Domains are not **strongly typed**; types are strongly typed

👉 Not all relational systems support user-defined domains

... “strongly typed” enforces strict restrictions on intermixing of values with differing data types

+ CREATE TABLE Example (3)

16

- A **referential triggered action** clause can be attached to a foreign key constraint, to specify the action to take if a referenced tuple is deleted, or a referenced primary key value is modified

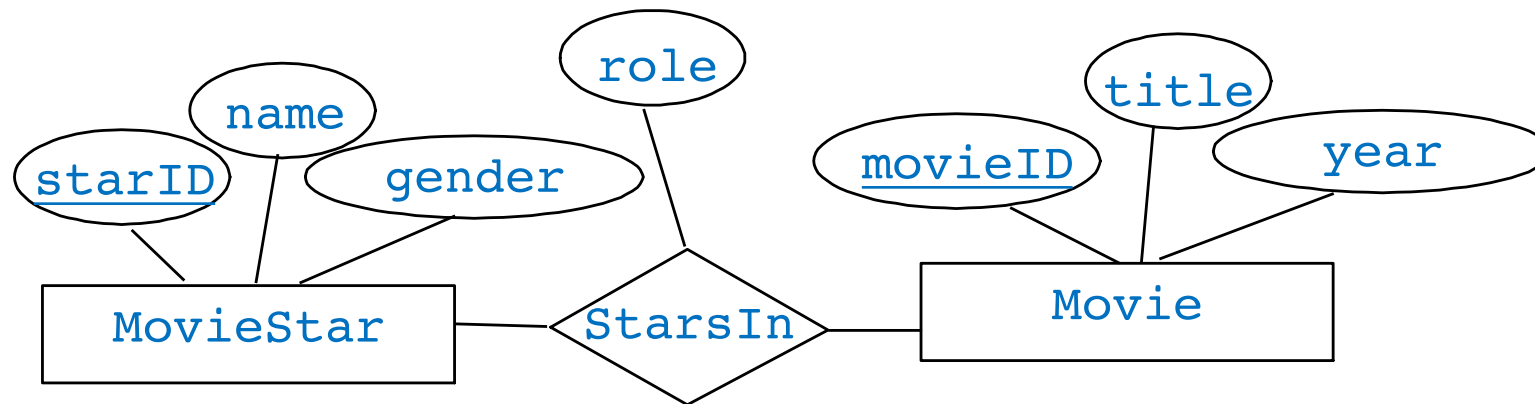
ON DELETE SET NULL | SET DEFAULT | CASCADE

ON UPDATE SET NULL | SET DEFAULT | CASCADE

```
CREATE TABLE Employee
(
    .....
    dNum          INT      NOT NULL DEFAULT 100,
    .....
    FOREIGN KEY (dNum) REFERENCES Department (dNumber)
        ON DELETE SET DEFAULT
        ON UPDATE CASCADE);
```


+ Another Example

17



`StarsIn[movieID, starID, role]`

`StarsIn.starID → MovieStar.starID`

`StarsIn.movieID → Movies.movieID`

```
CREATE TABLE StarsIn (  
  starID    INTEGER,  
  movieID   INTEGER,  
  role      CHAR(20),  
  PRIMARY KEY (starID, movieID),  
  FOREIGN KEY (starID) REFERENCES movieStar,  
  FOREIGN KEY (movieID) REFERENCES Movie);
```

+ Enforcing Referential Integrity

18

- `movieID` in `StarsIn` is a foreign key that references `Movie`
 - `StarsIn.movieID` → `Movie.movieID`
- What should be done if a `Movie` tuple is deleted, and there is a `StarsIn` tuple refers to it?
 1. Delete all roles that refer to it?
 2. Disallow the deletion of the movie?
 3. Set `movieID` in `StarsIn` tuples that refer to it to null?
 4. Set `movieID` in `StarsIn` tuples that refer to it to default value?

By default, no action is taken, and the delete/update is rejected.

Other actions include :

`ON DELETE SET NULL` | `SET DEFAULT` | `CASCADE`
`ON UPDATE SET NULL` | `SET DEFAULT` | `CASCADE`

+ Create Table with FK Actions

```
CREATE TABLE StarsIn (  
  starID    INTEGER,  
  movieID   INTEGER,  
  role      CHAR(20),  
  
  PRIMARY KEY (starID, movieID),  
  FOREIGN KEY (starID) REFERENCES MovieStar  
    ON DELETE CASCADE  
    ON UPDATE CASCADE,  
  FOREIGN KEY (movieID) REFERENCES Movie  
    ON DELETE SET NULL  
    ON UPDATE CASCADE);
```

+ Question:

20

Consider the following table definition.

```
CREATE TABLE ParkingPermit (  
    pID INT,  
    staffID INT, ...  
  
    PRIMARY KEY (pID),  
    FOREIGN KEY (staffID) REFERENCES Staff ON DELETE CASCADE);
```

Assume there is a tuple with `pID` = 1000 and `staffID` = 5678 in the table, choose the best answer

1. If the row for `staffID` value 5678 in `Staff` is deleted, then only the row with `pID` = 1000 in `ParkingPermit` is automatically deleted
2. If the row with `staffID` value 5678 in `Staff` is deleted, then all rows with `staffID`=5678 in `ParkingPermit` are automatically deleted
3. Both of the above

Based on the instruction given in the table definition, only option 2 is correct

+ ALTER TABLE

- **ALTER TABLE** command is used for schema evolution, that is the definition of a table created using the **CREATE TABLE** command, can be changed using the **ALTER TABLE** command
- Alter table actions include
 - Adding or dropping a column
 - Changing a column definition
 - Adding or dropping constraints

+ ALTER TABLE Syntax

22

ALTER TABLE <table name>

ADD <column name> <column type> [<attribute constraint>]

{, <column name> <column type> [<attribute constraint>] }

| DROP <column name> [CASCADE]

| ALTER <column name> <column-options>

| ADD <constraint name> <constraint-options>

| DROP <constraint name> [CASCADE];

👉 Commercial products have variations!

... to alter a constraint, it must be dropped and added again

... you can drop a PK constraint!

... FK doesn't have to reference to PK but can only to UNIQUE attributes; but in practice they always do

+ ALTER TABLE Examples

23

- To add an attribute

```
ALTER TABLE Employee ADD job VARCHAR(12);
```

- Note: values for the added attribute in all tuples will be initially NULL, so NOT NULL cannot be specified

- To drop an attribute

```
ALTER TABLE Employee DROP address;
```

- Note: drop an attribute which has been used by other tables in their FK references, **CASCADE** can be used

- To drop a constraint (constraint must have been given a name when it was specified)

```
ALTER TABLE Employee DROP CONSTRAINT  
empPK CASCADE;
```

+ Questions:

1. When adding a column, what happens to the existing records for their values of the added column?
2. When you drop a constraint, does it delete any data instances?
3. When you drop a PK constraint, do you have to specify cascade options?
4. When you drop an FK constraint, do you have to specify cascade options?
5. When two tables have mutual FKs, how do you insert data?

+ DROP TABLE

25

■ DROP TABLE

- Drops all constraints defined on the table including constraints in other tables which reference this table
- Deletes all tuples within the table
- Removes the table definition from the system catalog

■ DROP TABLE Syntax

DROP TABLE <table name> [**CASCADE**];

+ SQL Outline

26

- DDL

- Database Definition

- DML

- Database Modification

- SELECT Statements

- Database Views

- More on Database Constraints

+ Modifying Database Instances

27

- Statements to modify database **instances**
 - Those for modifying database **schemas** are called DDL
- Basic SQL Statements for modifying records
 - INSERT
 - DELETE
 - UPDATE

+ INSERT Statement

Used to add tuples to an existing relation

■ Single Tuple INSERT

- Specify the relation's name and a list of values for the tuple
- Values are listed in the same order as the attributes were specified in the CREATE TABLE command
- User may specify explicit attribute names that correspond to the values provided in the insert statement
 - The attributes not included cannot have the NOT NULL constraint

■ Multiple Tuple INSERT

- By separating each tuple's list of values with commas
- By loading the result of a query

+ INSERT Syntax

29

INSERT INTO <table name>

[(<column name> {, <column name> })]

VALUES (<constant value> {,<constant value> })

{(<constant value> {,<constant value> })}

| <select statement>;

...syntax check will be performed

+ INSERT Example: From Values

30

```
INSERT INTO Employee  
VALUES  
(  
    '653298695', 'Richard Marini', '30-DEC-1995',  
    '6B/98 University Road, Sai Kung, NT', 'M', 37000, 4, '987654321'  
);
```

```
INSERT INTO Employee(name, ssn)  
VALUES ('Jane Chow', '987654321'),  
       ('Emily Li', NULL);
```

Employee[ssn, name, dob, address, sex, salary, dNum, mgrSSN]

... syntax and integrity checks will be performed
... values with quotation or not can be flexible (i.e., supporting automatic
data type convention)

+ INSERT Example: From Queries

31

```
INSERT INTO DeptInfo(dName, numOfEmployees, totalSalary)
  SELECT      dName, COUNT(*), SUM(salary)
  FROM        Department, Employee
  WHERE       dNumber = dNum
  GROUP BY    dName ;
```

Employee[ssn,name,dob,address,sex,salary,dNum,mgrSSN]
Department[dName, dNumber, mgrSSN, mgrStartDate]

...we will discuss SELECT statement later

+ DELETE Statement

32

Used to remove some existing tuples from a relation

- Tuples are selected to delete from a **single** table
- Deletion may **propagate** to other tables if referential triggered actions are specified in the referential integrity constraints

```
DELETE FROM <table name>  
[WHERE <select condition>];
```

...DELETE all tuples doesn't equivalent to DROP a table

+ DELETE Example

33

```
DELETE FROM Employee  
WHERE dob < '1-JAN-1800';
```

```
DELETE FROM Employee  
WHERE dNum = 5;
```

```
DELETE FROM Employee  
WHERE salary >= 100000;
```

```
DELETE FROM Employee;
```

Employee[ssn, name, dob, address, sex, salary, dNum, mgrSSN]

... a single DELETE statement may delete zero, one, several or all tuples from a table

+ UPDATE Statement

Used to modify attribute values of one or more selected tuples in a relation

- Tuples are selected for update from a **single** table
- Updating a primary key value may **propagate** to other tables if referential triggered actions are specified in the referential integrity constraints

```
UPDATE <table name>  
SET <column name> = <value expression>  
  {, <column name> = <value expression>}  
[WHERE <select condition>];
```

+ UPDATE Example

35

```
UPDATE Employee  
SET      salary = salary * 1.1  
WHERE    lastName = 'McGowen';
```

Employee[ssn,name,dob,address,sex,salary,dNum,mgrSSN]

... a single UPDATE statement may change zero, one, several or all tuples from a table

+ More on DELETE/UPDATE

- Conceptually, deletion/update are done in two steps
 - Find the tuples to delete/update
 - Delete/update the tuples found

- Notes
 - They can only be used to delete/update records in one table, but can cause cascading changes in other tables
 - The where-clause can be as complex as in a SELECT statement, including using multiple relations
 - For UPDATE, only the values before the changes are considered in the find operation

+ SQL Outline

37

- DDL

- Database Definition

- DML

- Database Modification

- SELECT Statements

- Database Views

- More on Database Constraints

+ SELECT Statement

- SQL has one basic statement for retrieving information from the database
- In the SELECT statement, users specify what the result of the query should be, and the DBMS decides the operations and order of execution, thus SQL queries are **declarative**

+ SELECT Statements

- Simple SELECT queries
- Join queries
- Ordering your results
- Aggregation and Grouping
- Set operations
- Renaming

+ SELECT Basic Syntax

40

SELECT <attribute list>

FROM <table list>

[**WHERE** <condition>] ;

- <attribute list> is a list of attribute names whose values are to be retrieved by the query
- <table list> is a list of relation names required to process the query
- <condition> is a conditional (Boolean) expression that identifies the tuples to be retrieved by the query

+ Simple SELECT Example

41

```
SELECT      empAddress
FROM        Employee
WHERE       empName = 'Joe Bates';
```

Employee

empName	empAddress	department
Nicole Smith	1 Pine Road	CSE
Joe Bates	32 Chandler Rd	ECE

Query results

empAddress
32 Chandler Rd

+ SPJ Queries

- **Selection** (**WHERE** clause)


- Horizontally select tuples

- **Projection** (**SELECT** clause)

- Vertically select the attributes

- **Join** (**FROM** clause)

- Combine tuples from different relations for the search purposes



Conceptually, the Cartesian product of all the relations are generated, followed by applying the selection operation and then projection operation to produce the final results

+ Cartesian Product

43

- $R1 \times R2$: every row in $R1$ is combined with every row in $R2$ to form tuples in the result relation
- The schema of $R1 \times R2$ is the concatenation of all the columns from $R1$ and all the columns from $R2$

X	Y	Z	X			A	B	C
x1	y1	z1				a1	b1	c1
x2	y2	z2				a2	b2	c2

X	Y	Z	A	B	C
x1	y1	z1	a1	b1	c1
x1	y1	z1	a2	b2	c2
x2	y2	z2	a1	b2	c1
x2	y2	z2	a2	b2	c2

+ Selection in SQL

44

```
SELECT      <attribute list>  
FROM        <table list>  
[WHERE      {<join condition> AND} {<search condition>}];
```

- <search condition> is a conditional (Boolean) expression that identifies the tuples to be retrieved by the query
- <join condition> links tuples from the tables
 - Without any join condition, it means Cartesian Product
 - While joins can occur between any pair of attributes, most joins occur between primary keys and foreign keys

+ Selection Examples

45

List the names of employees working in department number 5.

```
SELECT  name
FROM    Employee
WHERE   dNum = 5;
```

List the names of employees who work in department 4 and earn over \$25000, or work in department 5 and earn over \$30000.

```
SELECT  name
FROM    Employee
WHERE   (dNum = 4 AND salary > 25000) OR
        (dNum = 5 AND salary > 30000);
```

Employee[ssn,name,dob,address,sex,salary,dNum,mgrSSN]

+ Join Query Example

46

List the names of employees working in the “Research” department.

```
SELECT    name
FROM      Employee, Department
WHERE     dNum = dNumber AND dName = 'Research';
```

```
Employee[ssn,name,dob,address,sex,salary,dNum,mgrSSN]
Department[dName, dNumber, mgrSSN, mgrStartDate]
```

+ Complex WHERE Conditions

■ Substring comparisons

■ LIKE

- WHERE Address LIKE '%Sai Kung%'
- WHERE mgrStartDate LIKE '__ / 0 5 / __'

■ IN

- WHERE lastName IN ('Jones', 'Wong', 'Chow')

■ IS

- WHERE dNum IS NULL

■ Arithmetic operators and functions

- +, -, *, /, date and time functions, etc.

- WHERE salary / 7.8 > 50000
- WHERE datediff(year, getdate(), dob) > 55

■ BETWEEN

- WHERE salary BETWEEN 10000 AND 30000

% : place holder for 0 or more characters
_ : place holder for a single character

+ Using Complex Conditions

48

List the names of employees in a research department with a salary between 40 and 60K.

```
SELECT    name
FROM      Employee, Department
WHERE     dNum = dNumber AND dName LIKE '%Research%'
          AND salary BETWEEN 40000 AND 60000;
```

```
Employee[ssn,name,dob,address,sex,salary,dNum,mgrSSN]
Department[dName, dNumber, mgrSSN, mgrStartDate]
```


+ Projection Examples

49

List the names, ssn and date of birth of all employees.

```
SELECT name, ssn, dob  
FROM Employee;
```

List all details of employees and their corresponding departments.

```
SELECT *  
FROM Employee, Department  
WHERE dNum = dNumber;
```

```
Employee[ssn, name, dob, address, sex, salary, dNum, mgrSSN]  
Department[dName, dNumber, mgrSSN, mgrStartDate]
```

+ Projection and Duplicates

50

- Theoretical, the projection operator produces a relation, which is a set of tuples
 - A set, by definition, cannot contain duplicates
- The SQL Projection can produce multisets (aka bags), which are tables with duplicate tuples
- Set semantics can be enforced in SQL using the **DISTINCT** option

+ Projection Example With DIST

51

List the salaries and departments of all employees.

```
SELECT salary, dDum  
FROM Employee;
```

```
SELECT DISTINCT salary, dDum  
FROM Employee;
```

salary	dNum
30000	5
40000	5
25000	4
43000	4
25000	5
25000	4
55000	1

Employee[ssn,name,dob,address,sex,salary,dNum,mgrSSN]

+ Question:

52

List the details of all employees.

```
SELECT * FROM Employee;
```

```
SELECT DIST * FROM Employee;
```

Will there be any difference in the results?

+ Projection and Expressions

53

- SQL queries can also evaluate expressions and return the value of these expressions together with the projected attributes
- Expressions use standard arithmetic operators (+, -, *, /) on numeric values or attributes with numeric domains, including a constant

+ Expression Example

54

List the names of all employees working in department 6, their salaries and salaries with a 17% increase.

```
SELECT    name, salary, 'Including super', 1.17 * salary
FROM      Employee
WHERE     dNum = 6;
```

```
SELECT    name, salary, 1.17 * salary AS 'includingSuper'
FROM      Employee
WHERE     dNum = 6;
```

Employee[ssn, name, dob, address, sex, salary, dNum, mgrSSN]

+ Sorting in SQL

55

```
SELECT [DISTINCT] <target list>
FROM   <table list>
[WHERE <condition>]
[ORDER BY <column>[ASC | DESC]
        {,<column>[ASC | DESC]} ] ;
```

target list: can be a column name, an expression, or *

column: can be a column name or a column position in the SELECT list

... sorting can be done by multiple columns

+ Sorting Examples

56

List all employee names and salaries, ordered by salary.

```
SELECT      name, salary
FROM        Employee
ORDER BY    salary DESC;
```

List all employee names, department number and salaries, ordered by department and salary.

```
SELECT      dNum, name, salary
FROM        Employee
ORDER BY    1 ASC, 3 DESC;
```

Employee[ssn,name,dob,address,sex,salary,dNum,mgrSSN]

+ What is Aggregation ?

57

- Aggregates are functions that produce summary values from a set of tuples
- Aggregates can be applied to
 - All tuples
 - A selected set of tuples
 - Multiple groups of tuples specified by the **GROUP BY** clause (to be discussed)
 - Aggregation functions can be used in the **SELECT** clause and the **HAVING** clause (to be discussed)

+ SQL Aggregation Functions

58

- SQL provides the following built-in functions for aggregates:
 - **COUNT**: Counts the number of tuples that the query returns
 - **SUM/AVG**: Calculates the sum (average) of a set of *numeric* values
 - **MAX/MIN**: Returns the maximum (minimum) value from a set of values which have a *total ordering*. Note that the domain of values can be non-numeric
 - They can be used with **DISTINCT**

+ Aggregation Examples

59

Find the total and average salary of employees.

```
SELECT    AVG(salary), SUM(salary)
FROM      Employee;
```

Find the age of the youngest employee in the
“Research” department.

```
SELECT    DATEDIFF(YEAR, MAX(dob), GetDate())
FROM      Employee, Department
WHERE     dNum = dNumber AND dName = 'Research';
```

```
Employee[ssn,name,dob,address,sex,salary,dNum,mgrSSN]
Department [dName, dNumber, mgrSSN, mgrStartDate]
```

+ More Aggregation Examples

60

Find the total number of employees in department 5.

```
SELECT      COUNT (*)  
FROM        Employee  
WHERE       dNum = 5;
```

Count the distinct salaries of employees in department 5.

```
SELECT      COUNT (DISTINCT salary)  
FROM        Employee  
WHERE       dNum = 5;
```

Note that COUNT (salary) will give the same result as COUNT (*), duplicate values are not eliminated

Employee[ssn,name,dob,address,sex,salary,dNum,mgrSSN]

+ Aggregation and Grouping

61

- An aggregation function can be applied to all rows of a table as previously shown
- However, aggregation functions are often applied to groups of rows within a table
- The **GROUP BY** clause provides this functionality

+ Grouping in SQL

62

```
SELECT [DISTINCT] <target list>  
FROM   <table list>  
[WHERE <condition>]  
[GROUP BY <grouping attributes> ]  
[HAVING   <group conditions> ]  
[ORDER BY column [ASC|DESC] {, column [ASC|DESC]}];
```

When GROUP BY is used in an SQL statement, any attribute appeared in SELECT Clause must appear either in GROUP BY clause or in an aggregation function!

+ Group By Examples

63

Find the total number of employees in each department.

```
SELECT      COUNT (*)  
FROM        Employee;
```

```
SELECT      dNum, COUNT (*)  
FROM        Employee  
GROUP BY    dNum;
```

```
SELECT      dNum, COUNT (*)  
FROM        Employee  
GROUP BY    dNum  
ORDER BY    2 ASC;
```

dNum	count
5	34
4	560
1	120
3	45

dNum	count
5	34
3	45
1	120
4	560

Employee[ssn,name,dob,address,sex,salary,dNum,mgrSSN]

+ More Group By Examples

64

Find the total number of employees with salary > 40000 in each department.

```
SELECT      dNum, COUNT (*)  
FROM        Employee  
WHERE       salary > 40000  
GROUP BY    dNum;
```

Find the total number of employees in each department, for each distinct salary.

```
SELECT      dNum, salary, COUNT (*)  
FROM        Employee  
GROUP BY    dNum, salary;
```

Employee[ssn,name,dob,address,sex,salary,dNum,mgrSSN]

+ Conditions on Groups

- Conditions can be imposed on the selection of groups to be included in the query result
- The **HAVING** clause (following the GROUP BY clause) is used to specify these conditions, similar to the WHERE clause
- Unlike the WHERE clause, the HAVING clause can also include aggregates

+ Group By Examples with Having

66

Find the total number of employees in departments with more than 5 employees.

```
SELECT      dNum, COUNT (*)  
FROM        Employee  
GROUP BY    dNum  
HAVING      COUNT (*) > 5;
```

Find the total number of employees who earn more than 40000, in departments with more than 5 employees.

```
SELECT      dNum, COUNT (*)  
FROM        Employee  
WHERE       salary > 40000  
GROUP BY    dNum  
HAVING      COUNT (*) > 5;
```

The query is wrong! It returns departments with more than 5 employees earning 40000. We will learn how to write the correct query later, which is non-trivial to write.

Employee[ssn,name,dob,address,sex,salary,dNum,mgrSSN]

+ Multiple Relation SQL Queries

67

- Basic Set Operations in SQL

- Union

- Intersection

- Difference/Minus

- Renaming in SQL

- Joins in SQL

+ Basic Set Operators

68

- Relation is a **set** of tuples (no duplicates; no order)

- Basic set operators apply to relations

- UNION

- INTERSECTION


- DIFFERENCE/MINUS

- Union compatibility



- The same number of columns

- Pair-wise compatible domains


Union of A and B



Intersection of A and B



Difference A minus B



Difference B minus A

$A \cup B$	commutative	$A \cup B = B \cup A$
	associative	$(A \cup B) \cup C = A \cup (B \cup C)$
$A \cap B$	commutative	$A \cap B = B \cap A$
	associative	$(A \cap B) \cap C = A \cap (B \cap C)$
$A - B$	<u>not</u> commutative	$A - B \neq B - A$
	<u>not</u> associative	$(A - B) - C \neq A - (B - C)$

+ Set Queries in SQL

69

```
SELECT eSSN FROM WorksOn  
UNION  
SELECT ssn FROM WorkedOn;
```

```
SELECT * FROM WorksOn  
UNION  
SELECT * FROM WorkedOn;
```

```
SELECT eSSN FROM WorksOn  
INTERSECT  
SELECT ssn FROM WorkedOn;
```

```
SELECT eSSN FROM WorksOn  
MINUS  
SELECT ssn FROM WorkedOn;
```

- **UNION ALL**: Results can contain duplicate tuples.
- A **UNION** effectively does `SELECT DISTINCT`
- You can do `(SELECT...) UNION (SELECT ...) ORDER BY...`
- This can also be applied to **Intersection** and **Minus**

`WorksOn[eSSN, pNo, hours]`

`WorkedOn[ssn, projNum, duration]`

+ Renaming in SQL

70

- SQL provides the facility to rename attribute and/or table names
 - Declaring an alias
- Renaming assists in shorten names, removing ambiguity and specifying self-joins

```
SELECT CustomerName,  
       CONCAT(address, ', ', country) AS fullAddress  
FROM   Customers;
```

```
SELECT emp.name, mgr.name AS MgrName  
FROM   Employee AS emp, Employee AS mgr  
WHERE  emp.mgrSSN = mgr.ssn;
```

.... Oracle can omit **AS** for table renaming

+ Joins in SQL

71

- A Join used to **combine related tuples** from two relations into a single tuple in a new (result) relation
- Join operation is needed for **organizing a search space** of data
 - This is needed when information is contained in more than one relation
- Join relations are specified in the **FROM Clause**.
When two relations are combined for a search, we need to know how the relations are combined
 - Based on **Cartesian Product** (denoted as X)
 - Denoted as ⋈

$$R_1 \bowtie \langle \text{join conditions} \rangle R_2$$

+ Cartesian Product in SQL

72

- Cartesian Product of two tables with no common attributes

```
SELECT *  
FROM Employee, Project;
```

```
Employee[ssn,name,dob,address,sex,salary,dNum,mgrSSN]  
Project[pNo, pName, pLocation, dNum]
```


+ Theta-Join

73

- A join operation is to select a subset of the Cartesian Product
- The most general type of join is called **θ -join**
- Join condition allows several logical operators $\{=, \neq, <, \leq, >, \geq\}$

$$\theta \in \{=, \neq, <, >, \leq, \geq\}$$

+ Equi-Join

- Joins tuples from two relations when they have the same value for some pair(s) of designated attributes
- The specification is termed a “join condition”
- With equi-join, the join condition only has **equality** comparisons

+ Join Examples

75

List the names of the managers of each department.

```
SELECT  E.name, D.dName
FROM    Department AS D, Employee AS E
WHERE   D.mgrSSN = E.ssn;
```

For each employee, list the employees who earn more than his/her manager.

```
SELECT  A.name, B.name
FROM    Employee A, Employee B
WHERE   A.salary > B.salary AND A.mgrSSN = B.ssn;
```

Employee[ssn,name,dob,address,sex,salary,dNum,mgrSSN]
Department [dName, dNumber, mgrSSN, mgrStartDate]

+ Natural Join

76

- Similar to equi-join except that the attributes that are used for the join are those that have **the same name** in each relation
 - Join attributes (and join condition) are not explicitly specified
 - Duplicate attributes are eliminated
 - No special SQL syntax for this join (some products can use keyword such as **NATURAL JOIN** in FROM-clause)
 - Natural join occurs frequently (due to table normalization)

+ Natural Join Example

77

emplID	dNum	dNum	head
1	101	101	Jane
2	210	210	Peter
3	101	300	Simon
4	300	400	Anna

emplID	dNum	head
1	101	Jane
2	210	Peter
3	101	Jane
4	300	Simon

+ Inner and Outer Joins

■ Inner Join

- This is the **default** join, in which a tuple is included in the result relation only if matching tuples exist in both relations

■ Outer Join

- Outer joins can include the tuples that do not satisfy the join condition, i.e, a matching tuple does not exist, which is indicated by a NULL value
- Full Outer Join: Includes all rows from both tables
- Left Outer Join: Includes all rows from first table
- Right Outer Join: Includes all rows from second table
- No special SQL syntax for this join (some products can use keyword such as **FULL OUTER JOIN** etc in FROM-clause)

+ Nested Queries in SQL

79

- Concept of Nested (sub) queries
- Correlated and non-correlated types
- Joins vs. Sub-queries
- Using the IN Operator
- Using the EXISTS Function
- Rules of sub-query construction

+ Nested SQL Queries

80

- A **nested query** (often termed **subquery**) is a query that appears within a query
 - Inside the WHERE clause of another SELECT statement
 - Inside an INSERT, UPDATE or DELETE statement
 - Nesting can occur at multiple levels
 - The query that contains the nested query is called its **outer query**
- Nested queries represent an alternative technique for expressing multi-table manipulation
- Useful for expressing queries where data must be fetched and used in a comparison condition

+ Two Types of Sub-queries

■ Non-correlated

- Results are returned from an inner query to an outer clause
- Sub-queries are evaluated from the “inside out”
- The outer query takes an action based on the results of the inner query

■ Correlated

- Conditions in subquery WHERE clause have references to some attributes of a relation in the outer query
- The outer SQL statement provides the values for the inner subquery to use in its evaluation
- The subquery is evaluated once for each (combination of) tuple in the outer query

+ Non-Correlated Query Example 1

82

Find the names of employees who work in a department managed by the employee with SSN = 11.

```
SELECT name
FROM Employee
WHERE dNum IN
      (SELECT dNumber
       FROM Department
       WHERE mgrSSN = 11);
```

Employee[ssn,name,dob,address,sex,salary,dNum,mgrSSN]
Department[dName, dNumber, mgrSSN, mgrStartDate]

+ Non-Correlated Query Example 2

83

Find the names of employees who work on projects in Dubai.

```
SELECT name
FROM Employee
WHERE ssn IN (SELECT ssn
               FROM WorksOn
               WHERE pNo IN
                     (SELECT pNum
                      FROM Project
                      WHERE pLocation = 'Dubai'));
```

Employee[ssn, name, dob, address, sex, salary, dNum, mgrSSN]

WorksOn[eSSN, pNo, hours]

Project[pNum, pName, pLocation, dNum]

+ Correlated Query Example

84

Find the names of departments that are responsible for a project located in Dubai.

```
SELECT    name
FROM      Department
WHERE     'Dubai' IN
           (SELECT    pLocation
            FROM      Project
            WHERE     dNum = dNumber);
```

```
Department[dName, dNumber, mgrSSN, mgrStartDate]
Project[pNum, pName, pLocation, dNum]
```

+ Naming in Correlated Queries

85

- There may be ambiguity, if attributes of outer and nested queries have the same name
- Reference to an unqualified attribute refers to the relation in the inner-most nested query
- If an attribute with the same name has to be referred from an outer query, it must be qualified by the relation (alias) name, e.g. `Department.dNum`

+ Correlated/Non-Correlated Queries

86

Find the name of departments which have projects in Dubai.

```
SELECT name
FROM   Department
WHERE  'Dubai' IN
      (SELECT pLocation
       FROM   Project
       WHERE  dNumber = dNum);
```

```
SELECT name
FROM   Department
WHERE  dNumber IN
      (SELECT dNum
       FROM   Project
       WHERE  pLocation = 'Dubai');
```

For $m=5$ departments and $n=100$ projects, there are $m \times n = 500$ scans in the correlated query but only $m+n=105$ scans in non-correlated query

```
Department[dName, dNumber, mgrSSN, mgrStartDate]
Project[pNum, pName, pLocation, dNum]
```

+ Joins vs Sub-queries

87

- Both joins and sub-queries are used for multi-table queries
 - They can often be used interchangeably
- Use sub-queries when you need to compare aggregates to other values
 - The aggregates can be generated from the nested query
- Use joins when you are displaying results from multiple tables
 - Results cannot be displayed from subqueries

+ Join Example

88

List department names for the departments that are located in the same place as a project.

```
SELECT  DISTINCT name
FROM    Department, Project
WHERE   dLoc = pLocation;
```

```
SELECT  DISTINCT name
FROM    Department D
WHERE   dLoc IN (SELECT  pLocation
                  FROM    Project);
```

```
Department[dName, dNumber, dLoc, mgrSSN, mgrStartDate]
Project[pNum, pName, pLocation, dNum]
```


+ Aggregation Example

89

List all employees with a salary equal to the lowest salary paid.

```
SELECT    MIN(salary)
FROM      Employee;
// Suppose Result = $15,500
```

```
SELECT    ssn, name, salary
FROM      Employee
WHERE     salary = 15500;
```

```
SELECT    ssn, name, salary
FROM      Employee
WHERE     salary = (SELECT MIN(salary)
                    FROM Employee);
```

Employee[ssn,name,dob,address,sex,salary,dNum,mgrSSN]

+ Sub-query Operators

90

A subquery S can return a set or a single value.

SET

- $\langle \text{exp} \rangle$ [NOT] IN (S)
 - exp is checked for set membership
- $\langle \text{exp} \rangle$ θ {ANY | ALL} (S)
 - exp is compared with the set returned
 - ANY: true for one tuple
 - ALL: true for all tuples

SINGLE VALUE

- $\langle \text{exp} \rangle$ θ (S)
 - exp is compared with the value returned
 - S must return a single value (otherwise an error will occur)

... $\theta = \{=, <, >, \geq, \leq, <>\}$

... Oracle uses SOME and ANY interchangeably

+ 'IN' And '= ANY' Are Equivalent

91

List department names for departments that have a manager with last name "Smith".

```
SELECT      DIST name
FROM        Department
WHERE       mgrSSN IN (SELECT      ssn
                        FROM        Employee
                        WHERE       name LIKE '%Smith');
```

```
SELECT      DISTINCT name
FROM        Department
WHERE       mgrSSN = ANY
                        (SELECT      ssn
                        FROM        Employee
                        WHERE       name LIKE '%Smith');
```

Employee[ssn,name,dob,address,sex,salary,dNum,mgrSSN]
Department[dName, dNumber, mgrSSN, mgrStartDate]

+ Comparison to ALL Example

92

List employees whose salary is greater than the salary of all employees in department 5.

```
SELECT *  
FROM Employee  
WHERE salary > ALL (SELECT salary  
FROM Employee  
WHERE dNum = 5);
```

Employee[ssn,name,dob,address,sex,salary,dNum,mgrSSN]
Department[dName, dNumber, mgrSSN, mgrStartDate]

+ Not IN and <>ANY

93

- IN and = ANY can be equivalent
- However, NOT IN and <> ANY are not!
- Instead, Not IN and <> ALL are equivalent
- Example
 - If a sub-query returns: {\$30K, \$32K, \$37K}
 - NOT IN means: NOT=\$30K AND NOT=\$32K AND NOT=\$37K
 - <> ANY means: NOT=\$30K OR NOT=\$32K OR NOT=\$37K
 - <> ALL means: NOT=\$30K AND NOT=\$32K AND NOT=\$37K

+ NOT IN Example

94

Find the names of departments that are not responsible for any project located in Perth.

```
SELECT    name FROM Department
WHERE     dNumber NOT IN
         (SELECT dNum FROM Project WHERE pLocation = 'Perth');
```

```
SELECT    name FROM Department
WHERE     dNumber <> ALL
         (SELECT dNum FROM Project WHERE pLocation = 'Perth');
```

```
SELECT    name FROM Department
WHERE     dNumber <> ANY
         (SELECT dNum FROM Project WHERE pLocation = 'Perth');
```

Department[dName, dNumber, dLoc, mgrSSN, mgrStartDate]
Project[pNum, pName, pLocation, dNum]

+ Correlated Subquery Example

95

Display employee name and department number for employees who earn a salary higher than the average in their department.

```
SELECT  E1.name, E1.dNum
FROM    Employee E1
WHERE   salary > (SELECT  avg(salary)
                  FROM    Employee E2
                  WHERE   E1.dNum = E2.dNum);
```

Employee[ssn, name, dob, address, sex, salary, dNum, mgrSSN]

+ Now Let's Look Back

96

Find the total number of employees who earn more than 40000, in departments with more than 5 employees.

```
SELECT      dNum, COUNT (*)
FROM        Employee
WHERE       salary > 40000
GROUP BY    dNum
HAVING      COUNT (*) > 5;
```

```
SELECT      dNum, Count(*)
FROM        Employee
WHERE       salary > 40000 AND
            dNum IN (SELECT dNum FROM Employee
                     GROUP BY dNum
                     HAVING   count(*) >5)
GROUP BY    dNum;
```

Employee[ssn,name,dob,address,sex,salary,dNum,mgrSSN]

+ Using the Exists Function

97

- The **EXISTS** function tests for the existence of data that meet the criteria of the sub-query
- **WHERE EXISTS (sub-query)**
 - Evaluates to true if the result of the correlated sub-query is a non-empty set, i.e. contains 1 or more tuples
 - Evaluates to true if the result of the correlated sub-query returns an empty set, i.e. zero tuples
- Subqueries are used with **EXISTS** and **NOT EXISTS** are **always correlated**

+ EXISTS Examples

98

Display the names of employees who have dependents.

```
SELECT  DIST name FROM Employee
WHERE   EXISTS (SELECT *
                  FROM   Dependent
                  WHERE  empSSN = ssn );
```

Display the names of employees who have no dependents.

```
SELECT  DIST name FROM Employee
WHERE   NOT EXISTS (SELECT *
                    FROM   Dependent
                    WHERE  empSSN = ssn );
```

Employee[ssn,name,dob,address,sex,salary,dNum,mgrSSN]
Dependent[empSSN,name,dob,relationship]

+ A Shophsicated Example

99

Display the details of employees who work on ALL projects.

```
SELECT * FROM Employee
WHERE NOT EXISTS
    (SELECT * FROM Project
     WHERE NOT EXISTS
        (SELECT * FROM WorksOn
         WHERE eSSN = Employee.ssn AND
               pNum = Project.pNum));
```

Employee[ssn,name,dob,address,sex,salary,dNum,mgrSSN]

WorksOn[eSSN,pNo,hours]

Project[pNum,pName,pLocation,dNum]

+ SQL Outline

■ DDL

- Database Definition

■ DML

- Database Modification
- SELECT Statements

■ Database Views

■ More on Database Constraints

+ Views

101

- Provide a way to hide certain data from certain users, or a convenience way for further queries
- Use DDL to create views
 - create view view-name as <query expression>
 - where: <query expression> is any legal SQL query

```
create view Branch_loan as  
select branch_name, loan_number  
from Loan;
```

```
select loan_number  
from Branch_loan  
where branch_name='Perryridge';
```

- Use DML to manipulate views
 - They were **virtual tables**, can be used as tables

`Loan[branch_name, loan_number, amount]`

+ View Update

102

- Assume that we allow users who have access to the `Branch_loan` view to insert records in the view
- Example: to add a new tuple to the `branch_loan` view.

```
insert into branch_loan values ('Perryridge', 'L-307');
```

This insertion must be represented by the insertion of the tuple ('Perryridge', 'L-307', *null*) into the `Loan` relation

`Loan[branch_name, loan_number, amount]`

`Branch_loan[branch_name, loan_number, amount]`

+ View Update Problems

103

- Suppose we want to insert the tuple ('Choi Hung', 'Lei Chen') into the `Branch_borrower` view

```
create view Branch_borrower as
select branch_name, customer_name
from Loan natural join Borrower;
```

- The `Loan` and `Borrower` relations have to be updated accordingly, violating integrity constraints!

```
insert into Loan (loan_number, amount, branch_name)
values (null, null, 'Choi Hung');
insert into Borrower (customer_name, loan_number)
values ('Lei Chen', null);
```

`Loan[branch_name, loan_number, amount]`
`Borrower[customer_name, loan_number]`

+ Rules For Updatable Views

104

- A view is updateable (i.e., tuples can be inserted, updated and deleted) if all of the following conditions are satisfied by the query that defines the view
 - The from clause has only one relation
 - The select clause contains only attribute names of the relation, and does not have any expressions, aggregates, or distinct specification
 - Any attribute not listed in the select clause can be set to null (i.e., it does not have a not null constraint and is not part of a primary key)
 - The query does not have a group by or having clause
- The rules for view update are often system dependent

+ SQL Outline

105

- DDL

- Database Definition

- DML

- Database Modification

- SELECT Statements

- Database Views

- More on Database Constraints

+ More on Integrity Constraints (IC)

106

- Domain constraints define **valid values** for attributes
- DBMS tests values inserted/updated to the database and check queries to ensure that the comparisons make sense
- The constraints that can be included:
 - **not null**: specifies that null values are not allowed
 - **primary key**: specifies a key for a relation (cannot be null)
 - **unique**: specifies that a set of attributes is a candidate key (can be null)
 - **foreign key**: specifies that one or more attributes refer to a primary key attribute in another relation
 - **check**: specifies a predicate that the values in every tuple must satisfy


+ Primary Key vs Unique Constraints

107

- A relation can possibly have many **candidate keys** but only one of them is chosen as the **primary key**

```
create table student (  
  student_id char (8),  
  name       char (20),  
  login      char (10),  
  age        int,  
  cga        real,  
  primary key (student_id)  
  unique (login));
```


```
create table enrolled (  
  student_id char (8),  
  course_id  char (10),  
  grade      char (2),  
  primary key (student_id),  
  unique (course_id, grade));
```



... used carelessly, an IC can prevent the storage of database instances that should be allowed!

+ Not Null & Participation Constraints

108



We can capture total participation constraints using **NOT NULL**

```
create table department (  
  did int,  
  dname char(20),  
  budget real,  
  HKID char(11) not null,  
  primary key (did),  
  foreign key (HKID) references employees on delete cascade);
```

+ CHECK CLAUSE: DOMAINS

109

- The check clause can be used in a create domain clause to add an integrity constraint to the domain.
- Example: ensure that an hourly-wage domain allows only values greater than a specified value (\$25 in this example)

```
create domain hourly_wage numeric(5, 2)  
constraint wage_test check (value >= 25.00);
```

- The new domain hourly_wage is declared to be a decimal number with 5 digits, 2 of which are after the decimal point
- The domain has a constraint named wage_test

👉 The constraint name is optional, but useful to indicate which constraint an update violated

+ Check Clause: Attributes

- The **check** clause can be used to add an integrity constraint for an attribute and can contain an arbitrary predicate
 - It is specified in the definition of a relation and checked whenever there is an update to the relation
 - Example: to ensure that semester can have only specified values and that the course id is between 1000 and 4999).

```
create table section (  
  course_id    char(8),  
  section_id   char(8),  
  semester     varchar(6),  
  year         numeric(4,0),  
  check (course_id between 1000 and 4999),  
  check (semester in ('Fall', 'Winter', 'Spring', 'Summer')));
```

+ Assertions

- An **assertion** is an arbitrary SQL predicate that the database must always satisfy

- An assertion takes the form:

```
create assertion <assertion-name> check <predicate>;
```

- Difference from previous constraints:
 - A constraint is associated with a single table and checked when there is an update on this specific table
 - An assertion may be associated with several tables and is checked every time when there is an update anywhere
- Assertion testing may introduce a significant amount of **overhead** and should be used with great care

+ Assertion Example

112

- The sum of all loan amounts for each branch must be less than the sum of all account balances at the branch

```
create assertion sum_constraint check (not exist (  
  select *  
  from branch  
  where (select sum(amount)  
         from loan  
         where loan.branch_name=branch.name)  
         >=  
         (select sum(amount)  
         from account  
         where account.branch_name=branch.name))));
```

- Since the assertion refers to multiple tables, it cannot be included as a constraint in the definition of `loan` or `account`

+ Triggers

113

- A **trigger** is a statement that is executed automatically by the system as a side effect of a modification to the database
- A trigger must
 - Specify the **condition**(s) under which it is to be executed
 - Specify the **action**(s) to be taken when it executes
- A trigger is used to implement integrity constraints that cannot be specified by SQL constraints
 - Example: automatically update the assets value of the branch relation whenever the balance value of the account relation is updated

+ Trigger Example

114

- Suppose that instead of allowing negative account balances, the bank deals with overdrafts by
 - Setting the account balance to zero
 - Creating a loan in the amount of the overdraft
 - Giving this loan a loan number which is identical to the account number of the overdrawn account
- The **condition** for executing the trigger is an update to the account relation that results in a negative balance value

+ Trigger Example

115

```
define trigger overdraft on update of account T
  (if new T.balance<0
   then (insert into loan values
         (T.branch_name, T.account_number, - new T.balance);
        insert into borrower
         (select customer_name, account_number
          from depositor
          where T.account_number=depositor.account_number);
        update account S
         set S.balance=0
         where S.account_number=T.account_number));
```

The keyword new used before T.balance indicates that the value of T.balance after the update should be used; if it is omitted, the value before the update is used. (T is the working version of account and S is the original version of account)

+ SQL Summary

- Structured Query Language (SQL) is a relational data manipulation language that provides facilities to
 - Query relations
 - Select-From-Where Statement
 - Set Operations (Union, Intersect, Except)
 - Nested Subqueries (to test for set membership, comparison, cardinality)
 - Aggregate Functions (avg, min, max, sum, count)
 - Group By with Having clause
 - Create and modify relations
 - Create, Alter, Drop Tables
 - Specify integrity constraints: domain, key, foreign key, general
 - Specify views
 - Insert, Delete, Update Tuples

+ Readings

- Textbook: Ch 3, 4, 5 (ed 6/ed 7)

Further References:

- W3School SQL tutorial
 - <https://www.w3schools.com/sql/>
- Oracle SQL
 - https://docs.oracle.com/cd/B19306_01/server.102/b14200/toc.htm
- Oracle University
 - <https://education.oracle.com/>